

Pricing

IndividualTeam & EnterpriseAPI

Latest models

Claude Opus 4

- Most intelligent model for complex tasks
- 200K context window
- 50% discount with batch processing*

Input
\$15 / MTok

Prompt caching write
\$18.75 / MTok

Output
\$75 / MTok

Prompt caching read
\$1.50 / MTok

Claude Sonnet 4

- Optimal balance of intelligence, cost, and speed
- 200K context window
- 50% discount with batch processing*

Input
\$3 / MTok

Prompt caching write
\$3.75 / MTok

Output
\$15 / MTok

Prompt caching read
\$0.30 / MTok

Claude Haiku 3.5

- Fastest, most cost-effective model**
- 200K context window
- 50% discount with batch processing*

Input
\$0.80 / MTok

Prompt caching write
\$1 / MTok

Output
\$4 / MTok

Prompt caching read
\$0.08 / MTok

*Learn more about [batch processing](#)

**[Claude Haiku 3.5](#) in Amazon Bedrock supports latency optimization for 60% faster inference speed, starting at \$1 per million input tokens and \$5 per million output tokens

Customers can purchase prioritized API capacity with [Priority Tier](#)

Prompt caching pricing is for our standard 5-minute TTL; [extended prompt caching](#) is available at an additional cost

Start building

Explore pricing for tools

Web search

- Give Claude access to the latest information from the web

Cost
\$10 / 1K searches*

Code execution

- Run Python code in a sandboxed environment for advanced data analysis
- 50 free hours of usage daily per organization

Cost
\$0.05 per hour per container
for additional usage

*Does not include input and output tokens required to process requests

Start building

Explore legacy models

Claude Opus 3

- 200K context window
- 50% discount with batch processing*

Input
\$15 / MTok

Prompt caching write
\$18.75 / MTok

Output
\$75 / MTok

Prompt caching read
\$1.50 / MTok

Claude Sonnet 3.7

- 200K context window
- 50% discount with batch processing*

Input
\$3 / MTok

Prompt caching write
\$3.75 / MTok

Output
\$15 / MTok

Prompt caching read
\$0.30 / MTok

Claude Haiku 3

- 200K context window
- 50% discount with batch processing*

Input
\$0.25 / MTok

Prompt caching write
\$0.30 / MTok

Output
\$1.25 / MTok

Prompt caching read
\$0.03 / MTok

Prompt caching pricing is for our standard 5-minute TTL; [extended prompt caching](#) is available at an additional cost

Learn more



Product

- Claude overview
- Claude Code
- Claude team plan
- Claude enterprise plan
- Claude education plan
- Download Claude apps
- Claude.ai pricing plans
- Claude.ai login

API Platform

- API overview
- Developer docs
- Claude in Amazon Bedrock
- Claude on Google Cloud's Vertex AI
- Pricing
- Console login

Research

- Research overview
- Economic Index

Claude models

- Claude Opus 4
- Claude Sonnet 4
- Claude Haiku 3.5

Commitments

- Transparency
- Responsible scaling policy
- Security and compliance

Solutions

- AI agents
- Coding
- Customer support

Learn

- Anthropic Academy
- Customer stories
- Engineering at Anthropic
- Explore
- About us
- Become a partner
- Careers
- Events
- News
- Startups program

Help and security

- Status
- Availability
- Support center
- Terms and policies
- Privacy choices
- Privacy policy
- Responsible disclosure policy
- Terms of service - consumer
- Terms of service - commercial
- Usage policy

Cookie Settings

We use cookies to deliver and improve our services, analyze site usage, and if you agree, to customize or personalize your experience and market our services to you. You can read our [Cookie Policy](#) [here](#).

Customize cookie settings

Reject all cookies

Accept all cookies